



# AE 03: Joining prognosticators

## Suggested answers

APPLICATION EXERCISE

ANSWERS

MODIFIED

September 13, 2024

### Important

These are suggested answers. This document should be used as reference only, it's not designed to be an exhaustive key.

```
library(tidyverse)
library(scales)

seers <- read_csv("data/prognosticators.csv")
weather <- read_csv("data/weather-region.csv")
```

## Prognosticator success

We previously examined the accuracy rate of Groundhog Day prognosticators.<sup>1</sup> Today we want to work with the original dataset to understand how those accuracy metrics were generated and answer the question:

**How does prognosticator accuracy vary by climatic region?**

<sup>1</sup> See [ae-01](#)

Let's start by looking at the `seers` data frame.

```
glimpse(seers)
```

Rows: 1,710

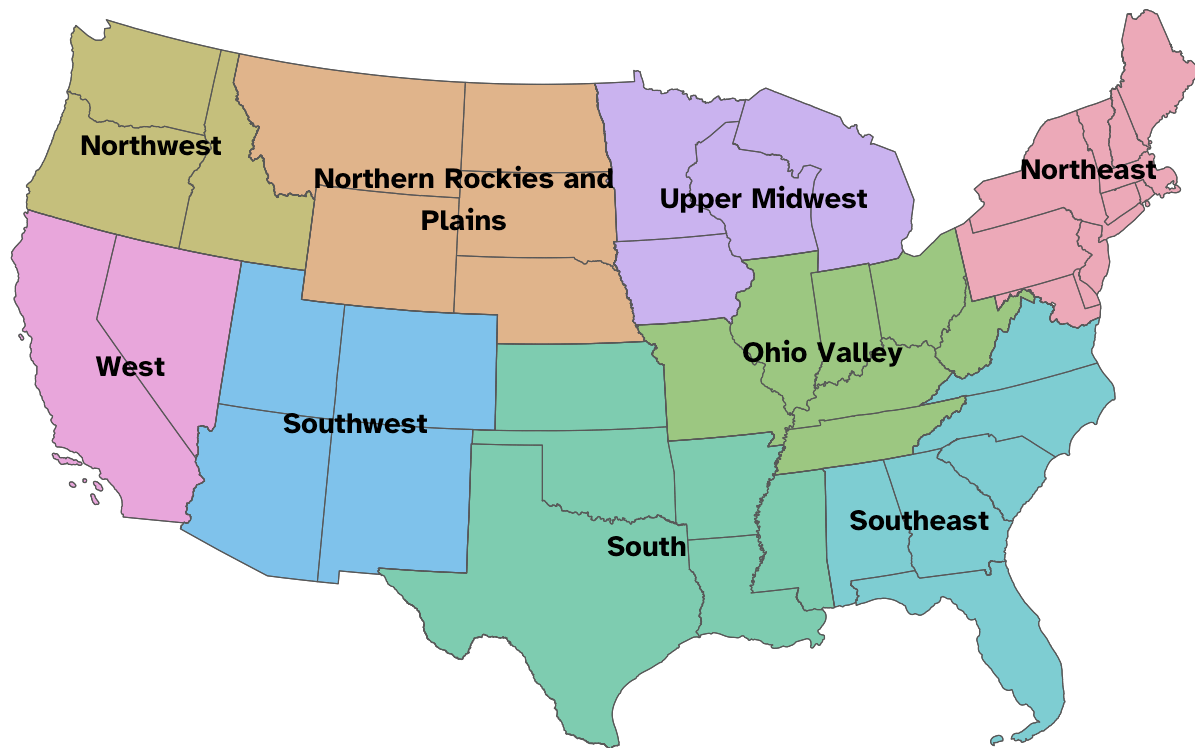
Columns: 7

```
$ name           <chr> "Punxsutawney Phil", "Punxsutawney Phil", "Punxsutawne...
$ forecaster_type <chr> "Groundhog", "Groundhog", "Groundhog", "Groundhog", "G...
$ alive          <lgl> TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, ...
$ town           <chr> "Punxsutawney", "Punxsutawney", "Punxsutawney", "Punxs...
$ state          <chr> "PA", "PA", "PA", "PA", "PA", "PA", "PA", "PA", "PA", ...
$ year           <dbl> 2024, 2023, 2022, 2021, 2020, 2019, 2018, 2017, 2016, ...
$ prediction     <chr> "Early Spring", "Late Winter", "Late Winter", "Late Wi...
```

We have the predictions, but our goal is to make a visualization by climate region.<sup>2</sup>

<sup>2</sup> Source: [National Weather Service Climate Prediction Center](#)

## Climatic regions in the United States



Source: National Climatic Data Center

## Join the data frames

Let's take a look at the weather data frame.

```
glimpse(weather)
```

Rows: 5,568

Columns: 13

```
$ region      <chr> "Northeast", "Northeast", "Northeast", "Northeast", "No...
$ state_abb   <chr> "CT", "CT", "CT", "CT", "CT", "CT", "CT", "CT", "CT", "..."
$ id          <dbl> 101, 101, 101, 101, 101, 101, 101, 101, 101, 101, 101, ...
$ year        <dbl> 1909, 1910, 1911, 1912, 1913, 1914, 1915, 1916, 1917, 1...
$ avg_temp    <dbl> 28.00, 29.20, 24.90, 23.15, 28.05, 22.05, 27.50, 21.55,...
$ temp_hist   <dbl> 25.58333, 26.09000, 26.16667, 25.85667, 25.63333, 25.52...
$ temp_hist_sd <dbl> 4.245360, 4.241218, 4.103158, 4.124311, 3.907804, 4.016...
$ temp_sd     <dbl> 4.154767, 4.154767, 4.154767, 4.154767, 4.154767, 4.154...
$ precip      <dbl> 4.005, 2.520, 2.810, 3.570, 3.765, 2.920, 2.330, 3.425,...
$ precip_hist <dbl> 3.476667, 3.526667, 3.378000, 3.411000, 3.446333, 3.352...
$ precip_hist_sd <dbl> 1.1784719, 1.2081292, 1.1442431, 1.1620681, 1.2039309, ...
```

```
$ precip_sd      <dbl> 0.9715631, 0.9715631, 0.9715631, 0.9715631, 0.9715631, ...
$ outcome        <chr> "Early Spring", "Early Spring", "Early Spring", "Late W...
```

- **Your turn (2 minutes):**

- Which variable(s) will we use to join the `seers` and `weather` data frames?
- We want to keep all rows and columns from `seers` and add columns for corresponding weather data. Which join function should we use?

- **Demo:** Join the two data frames and assign the joined data frame to `seers_weather`.

```
seers_weather <- inner_join(
  x = seers, y = weather,
  by = join_by(state == state_abb, year)
)
```

## Calculate the variables

- **Demo:** Take a look at the updated `seers` data frame. First we need to calculate for each prediction whether or not the prognostication was correct.

```
seers_weather <- seers_weather |>
  mutate(correct_pred = prediction == outcome)
```

- **Demo:** Calculate the accuracy rate (we'll call it `preds_rate`) for weather predictions using the `summarize()` function in `dplyr`. Note that the function for calculating the mean is `mean()` in R.

```
seers_weather |> # start with seers data frame
  group_by(region) |> # group by region
  summarize(preds_rate = mean(correct_pred)) # calculate accuracy rate
```

# A tibble: 9 × 2

	region	preds_rate
	<chr>	<dbl>
1	Northeast	0.491
2	Northern Rockies and Plains	0.574
3	Northwest	0.442
4	Ohio Valley	0.557
5	South	0.506
6	Southeast	0.568
7	Southwest	0.667
8	Upper Midwest	0.5
9	West	0.286

- **Your turn (5 minutes):** Now expand your calculations to also calculate the number of predictions in each region and the standard error of accuracy rate. Store this data frame as `seers_summary`. Recall the formula for the standard error of a sample proportion:

$$SE(\hat{p}) \approx \sqrt{\frac{(\hat{p})(1 - \hat{p})}{n}}$$

```
seers_summary <- seers_weather |>
  group_by(region) |>
  summarize(
    preds_rate = mean(correct_pred),
    preds_n = n(),
    preds_se = sqrt((preds_rate * (1 - preds_rate)) / preds_n)
  )
seers_summary
```

# A tibble: 9 × 4

	region	preds_rate	preds_n	preds_se
	<chr>	<dbl>	<int>	<dbl>
1	Northeast	0.491	696	0.0189
2	Northern Rockies and Plains	0.574	47	0.0721
3	Northwest	0.442	52	0.0689
4	Ohio Valley	0.557	280	0.0297
5	South	0.506	79	0.0562
6	Southeast	0.568	176	0.0373
7	Southwest	0.667	36	0.0786
8	Upper Midwest	0.5	104	0.0490
9	West	0.286	7	0.171

- **Demo:** Take the `seers_summary` data frame and order the results in descending order of accuracy rate.

```
seers_summary |> # start with seers_summary data frame
  arrange(desc(preds_rate)) # order in descending order of preds_rate
```

# A tibble: 9 × 4


	region	preds_rate	preds_n	preds_se
	<chr>	<dbl>	<int>	<dbl>
1	Southwest	0.667	36	0.0786
2	Northern Rockies and Plains	0.574	47	0.0721
3	Southeast	0.568	176	0.0373
4	Ohio Valley	0.557	280	0.0297
5	South	0.506	79	0.0562
6	Upper Midwest	0.5	104	0.0490
7	Northeast	0.491	696	0.0189
8	Northwest	0.442	52	0.0689
9	West	0.286	7	0.171

## Recreate the plot

- **Demo:** Recreate the following plot using the data frame you have developed so far.

## Prognosticator accuracy rate for late winter/early spring

By climate region

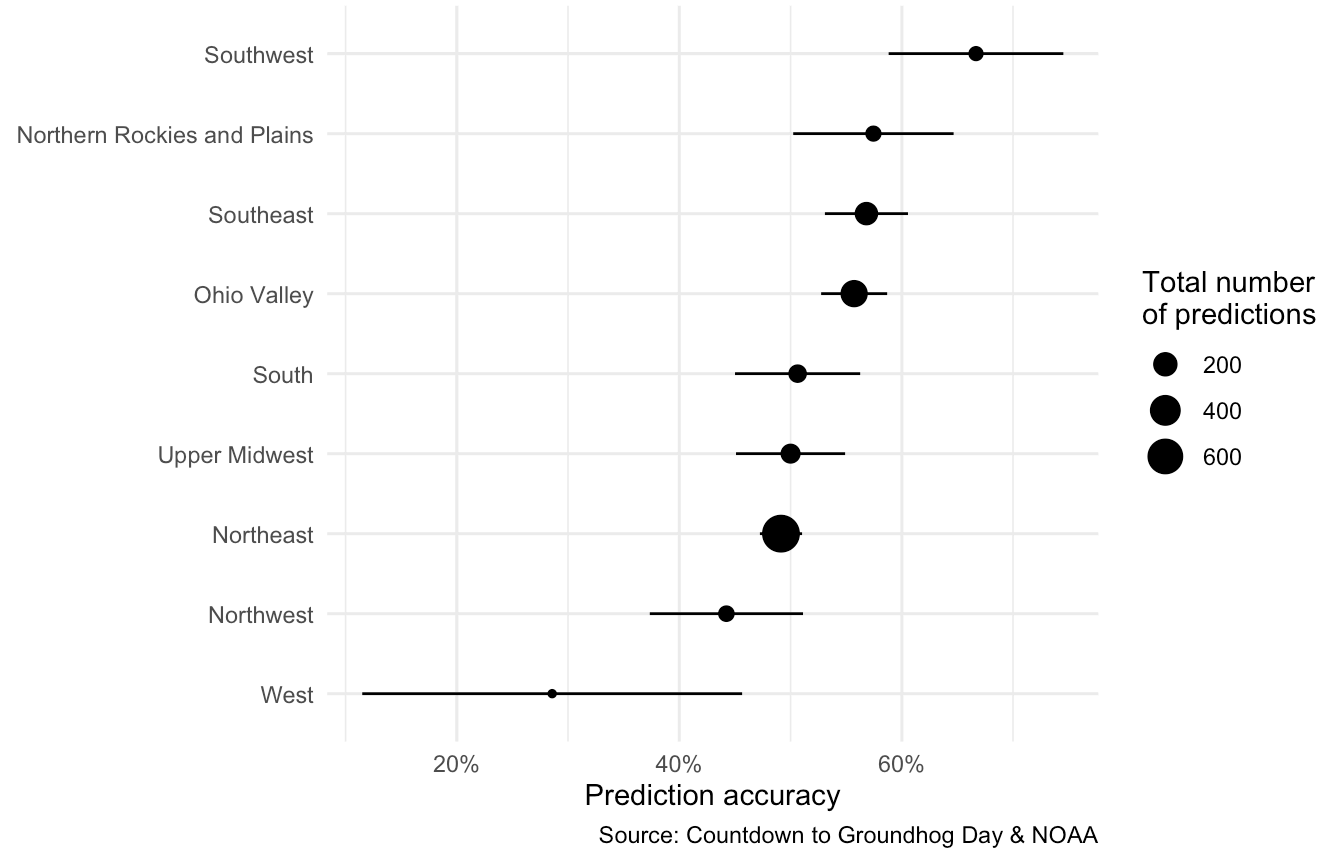


Source: Countdown to Groundhog Day & NOAA

```
seers_summary |>
  mutate(region = fct_reorder(.f = region, .x = preds_rate)) |>
  ggplot(mapping = aes(x = preds_rate, y = region)) +
  geom_point(mapping = aes(size = preds_n)) +
  geom_linerange(mapping = aes(
    xmin = preds_rate - preds_se,
    xmax = preds_rate + preds_se
  )) +
  scale_x_continuous(labels = label_percent()) +
  labs(
    title = "Prognosticator accuracy rate for late winter/early spring",
    subtitle = "By climate region",
    x = "Prediction accuracy",
    y = NULL,
    size = "Total number\nof predictions",
    caption = "Source: Countdown to Groundhog Day & NOAA"
  ) +
  theme_minimal()
```

Prognosticator accuracy rate for late winter/early spring

By climate region



- **Your turn (time permitting):** Make any other changes you would like to improve it.

# add your code here

Session information